



**ML6**

DATA SCIENCE  
GHENT



OZARK  HENRY  
**boobook**



# Agenda

- Today 19.30: explanation of problem, dataset & forming teams
- Today 20.15: official start hackathon!
- Today  21.30: pizza's
- Tomorrow  09.00: start breakfast @digityser
- Tomorrow  14.30: start lunch
- Tomorrow  20.30: start dinner
- Sunday  09.00: start brunch
- Sunday 12.30: end of hackathon (deadline for submissions)
- Sunday 13.00: start of award ceremony with guest: Ozark Henry!
- Sunday 15.00: end of ceremony & reception

# Tweettweet tweet



#MusicHack

@DSGhent

@ML6team

@DigitYser

# Practicalities

- Every participant has to sign the general terms and conditions sheet @ reception
- Please respect the premises & other people's stuff
- Smoking is permitted on the roof terrace
- Chill zone in the basement
- Wifi: DIGITYSER-2.4 or 5, password: Easy24Get
- Our phone numbers:



Matthias

+32 498/11.83.09



Sam

+32 485/58.54.52



Xander

+32 473/86.74.89



Hendrik

+32 484/29.50.09

# Food

- Food: Register on [Eventbrite](#) and pay €30 for all-in food package during the weekend, or bring your own food

I need no Invoice ticket - pay 30€ cash at  
the bar  
FREE

0

- Water and coffee are free, any other drink can be taken from the HONESTY bar (pay when you drink)


# Hackathon Flow

1. Form team and register on: [link to spreadsheet](#)
2. Hack your way towards an awesome, NN-generated MP3 file!
3. Send us your solution (before Sunday 12.30)
  - MP3 (max 3) upload to: [google form](#)
  - Add your slide to slidedeck: [final presentation](#)
    - Up to 3 solutions allowed per team: ~5 slides per solution

The background is an abstract, fluid composition of organic, wavy shapes. On the left, there's a large, vibrant green area with yellow highlights, resembling a cross-section of a tree trunk or a topographical map. This transitions into a bright yellow, glowing band that curves across the middle. To the right, the colors shift into deep blues and purples, with some lighter, misty areas. The overall effect is one of complex, layered textures and dynamic movement.

Music is complicated

# Raw waveform





Pressure


Time


D294

A440

F349

C523







$$f(t) = A_0 + \sum_{n=1}^{\infty} (A_n \cos(nt) + B_n \sin(nt))$$



# Spectrum Analyze



# Spectrogram







# DeepMind WaveNet




1 Second





# DeepMind WaveNet



# MIDI

Musical Instrument Digital Interface



# Get started with MIDI

## Lakh MIDI Dataset Tutorial

This IPython notebook demonstrates how to use the data in the [Lakh MIDI Dataset](#). It shows how the dataset is organized and gives examples of how to use annotations extracted from LMD-aligned (the collection of MIDI files which have been matched and aligned to entries in the Million Song Dataset). We will use [pretty\\_midi](#) for parsing the MIDI files, [mir\\_eval](#) for sonification and visualization, and [librosa](#) for audio analysis.


```
In [1]: # Imports
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import pretty_midi
import librosa
import mir_eval
import mir_eval.display
import tables
import IPython.display
import os
import json

# Local path constants
DATA_PATH = 'data'
RESULTS_PATH = 'results'
# Path to the file match_scores.json distributed with the LMD
SCORE_FILE = os.path.join(RESULTS_PATH, 'match_scores.json')
```

<http://bit.ly/2FAHuv3>

```
# Utility functions for retrieving paths
def msd_id_to_dirs(msd_id):
    """Given an MSD ID, generate the path prefix
```





# magenta

<https://magenta.tensorflow.org/>



# Start your own Linux VM on gcp

## Compute Engine

[Product Overview](#)[Documentation](#)

## Quickstarts

[All Quickstarts](#)[Using a Linux VM](#)[Using a Windows VM](#)

## How-to Guides

[All How-to Guides](#)[▶ Creating VM Instances](#)[▶ Connecting to Instances](#)[▶ Adding Storage](#)[▶ Creating and Managing Instance](#)[▶ Templates](#)[▶ Creating and Managing Custom](#)[▶ Images](#)[▶ Managing Your Instances](#)[▶ Creating and Managing Groups of](#)[▶ Instances](#)[▶ Networking](#)[▶ Deploying Containers](#)[Compute Engine](#) > [Documentation](#)[FEEDBACK VERZENDEN](#)

## Quickstart Using a Linux VM

This page explains how to create a Linux virtual machine instance in Compute Engine using the Google Cloud Platform Console.

### Before you begin

1. Select or create a Cloud Platform project.

[GO TO THE MANAGE RESOURCES PAGE](#)

2. Factureren voor uw project inschakelen.

[FACTURERING INSCHAKELEN](#)

### Create a virtual machine instance

1. In the GCP Console, go to the VM Instances page.

[GO TO THE VM INSTANCES PAGE](#)

# Install Magenta Development Version

- `sudo apt-get update`
- `sudo apt-get install libasound2-dev libjack-dev python-pip htop`
- `pip install --upgrade pip`
- `sudo pip install matplotlib scipy bokeh IPython pandas jupyter`
- `git clone https://github.com/tensorflow/magenta.git`

## ***Install Bazel:***

- `sudo apt-get install openjdk-8-jdk`
- `echo "deb [arch=amd64] http://storage.googleapis.com/bazel-apt stable jdk1.8" | sudo tee /etc/apt/sources.list.d/bazel.list`  
`curl https://bazel.build/bazel-release.pub.gpg | sudo apt-key add -`
- `sudo apt-get update && sudo apt-get install bazel`
- `sudo apt-get upgrade bazel`
- `sudo pip install magenta`

# Generate your first MIDI-song:

## Download pretrained RNN:

- `cd magenta/magenta/models/melody_rnn`
- `mkdir trained_models`
- `mkdir generated_midis`
- `curl -o trained_models/attention_rnn.mag http://download.magenta.tensorflow.org/models/attention\_rnn.mag`

## Generate MIDI using attention\_rnn: *(replace “.” with the correct directories)*

```
python melody_rnn_generate.py \  
--config='attention_rnn' \  
--bundle_file= /home/./magenta/magenta/models/melody_rnn/trained_models/attention_rnn.mag \  
--output_dir=generated_midis \  
--num_outputs=10 \  
--num_steps=128 \  
--primer_melody="[60]"
```

Generate MIDI to mp3: <http://mp3-tools.com/free-midi-to-mp3-converter.html>

# MIDI Synthesiser

Magenta has a custom MIDI synthesiser that allows you to play with various types of MIDI-to-sound scripts:

<https://github.com/tensorflow/magenta/tree/master/magenta/interfaces/midi>



EXPLORE **SOUND**

# Get + explore MIDI data: (includes very useful notebooks!)

Explore the **‘Lakh MIDI Dataset v0.1’**:

<http://colinraffel.com/projects/lmd/>

**Download and store the MIDI dataset on the VM:**

```
mkdir midi_data
```

```
curl -o midi_data/lmd_aligned.tar.gz http://hog.ee.columbia.edu/craffel/lmd/lmd\_aligned.tar.gz
```

```
curl -o midi_data/clean_midi.tar.gz http://hog.ee.columbia.edu/craffel/lmd/clean\_midi.tar.gz
```

```
cd midi_data
```

```
tar -xvzf lmd_aligned.tar.gz
```

```
tar -xvzf clean_midi.tar.gz
```



Download TfreCORDS files: (avoid 4 hours of waiting)

[clean\\_midi.tfreCORDS](#)

[sequence\\_data.zip](#)

# Generate TF\_Records file from MIDI-data:

```
cd ../scripts/
```

```
INPUT_DIRECTORY=/home/../../midi_data/clean_midi/
```

```
SEQUENCES_TFRECORD=/home/../../midi_data/clean_midi.tfrecord
```

```
#Run script headless (takes a really long time!!)
```

```
nohup python convert_dir_to_note_sequences.py \
```

```
--input_dir=$INPUT_DIRECTORY \
```

```
--output_file=$SEQUENCES_TFRECORD \
```

```
--recursive &> clean_midi_log.out&
```

**To check that your script is running you can type “htop” in the command line**

**The output of the script will be written to “clean\_midi\_log.out”**



# Create Train/Validation/Test sets

## Go back to root folder

```
mkdir midi_data/sequence_data
```

```
cd ../../models/melody_rnn
```

```
nohup python melody_rnn_create_dataset.py \  
--config='attention_rnn' \  
--input=/home/../../midi_data/clean_midi.tfrecord \  
--output_dir=/home/../../midi_data/sequence_data \  
--eval_ratio=0.10 &> generate_train_data.out&
```

A high-angle, nighttime photograph of a massive crowd of people gathered in an open space. The crowd is densely packed, filling most of the frame. In the background, a city skyline is visible with numerous lights reflecting on a body of water. The overall scene suggests a large-scale event or festival.

**Are  
You  
Ready?**

# Create Train/Validation/Test sets

```
python melody_rnn_train.py \  
--config=attention_rnn \  
--run_dir=/tmp/melody_rnn/logdir/run1 \  
--sequence_example_file=/home/../../midi_data/sequence_data/  
training_melodies.tfrecord \  
--hparams="batch_size=64,rnn_layer_sizes=[64,64]" \  
--num_training_steps=20000
```

# How to register your team

Go to


[https://docs.google.com/spreadsheets/d/1Tdgm-vsX8hMVRVjfD9FLSs3kuQNwuC  
KITdcA8pbHDt4/edit#gid=0](https://docs.google.com/spreadsheets/d/1Tdgm-vsX8hMVRVjfD9FLSs3kuQNwuCKITdcA8pbHDt4/edit#gid=0)

and register your team and its members


# Attention RNN:


```
DEFAULT_MIN_NOTE = 48
DEFAULT_MAX_NOTE = 84
DEFAULT_TRANSPOSE_TO_KEY = 0

'attention_rnn': MelodyRnnConfig(
    magenta.protobuf.generator_pb2.GeneratorDetails(
        id='attention_rnn',
        description='Melody RNN with lookback encoding and attention.'),
    magenta.music.KeyMelodyEncoderDecoder(
        min_note=DEFAULT_MIN_NOTE,
        max_note=DEFAULT_MAX_NOTE),
    tf.contrib.training.HParams(
        batch_size=128,
        rnn_layer_sizes=[128, 128],
        dropout_keep_prob=0.5,
        attn_length=40,
        clip_norm=3,
        learning_rate=0.001))
```




=





The repeating module in an LSTM contains four interacting layers.

Don't worry about the details of what's going on. We'll walk through the LSTM diagram step by step later. For now, let's just try to get comfortable with the notation we'll be using.



# Interesting Links:

- [Music Generation Challenge on CrowdAI](#)
- [More info on the MIDI file format](#)
- [Other GitHub Repo's for music generation](#)
- [How LSTM's work](#) (Colah's Blog)
- 2 Episodes from Siraj' on music generation:
  - [Generate Music in TensorFlow](#)
  - [How to Generate Music with Tensorflow \(LIVE\)](#)





## Slightly related Links *(if you need some chill-out)*

- [Rising pitch sound effect in Dunkirk](#)



# Useful Linux commands:

**Run Jupyter notebook @VM: (paste this in SSH terminal and copy the login token)**

```
jupyter notebook --ip=0.0.0.0 --port=8888 --no-browser & disown
```

**Run Python script headless & append to log.out:**

```
sudo nohup python run.py &> log.out&
```

**List (& kill) all running jupyter servers:**

```
jupyter notebook list
```

```
kill $(pgrep jupyter)
```

**Check running python jobs:**

```
pgrep -af python
```

**Check disk & CPU usage:**

```
htop
```

```
df
```

# Evaluation

1. MP3 + Technical presentation
2. Public voting (1 - 2 weeks after Hackathon)



Many thanks to

**Link to this slidedeck:**  
**[goo.gl/8Gt87S](https://goo.gl/8Gt87S)**

